

High-Performance Web-Services in C++

C++ User Group Dresden
Maximilian Haupt
2016/09/08

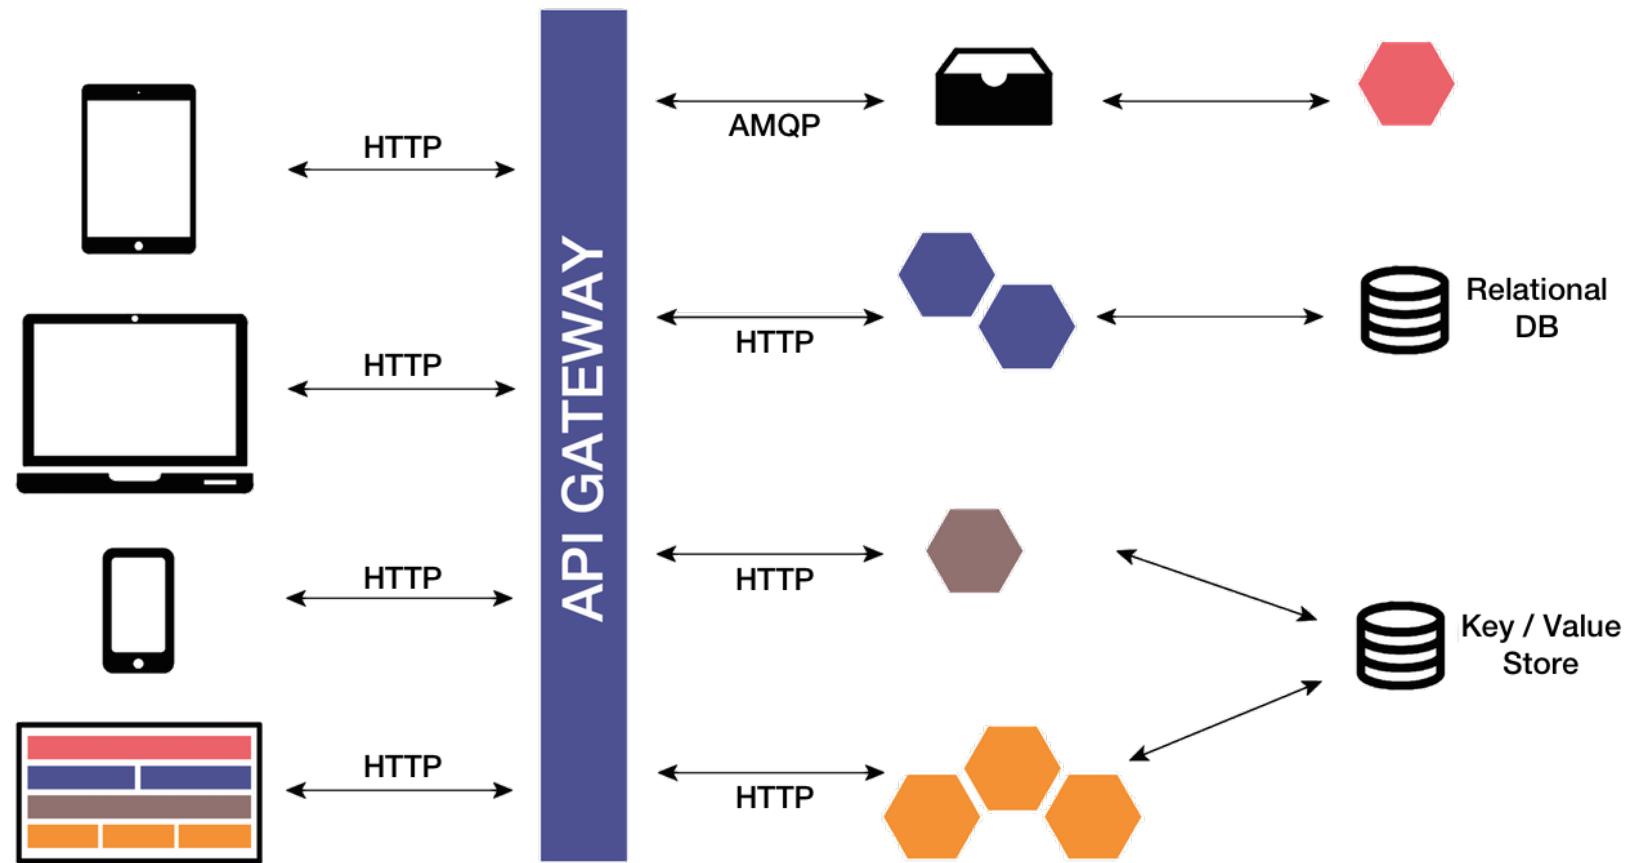
Agenda

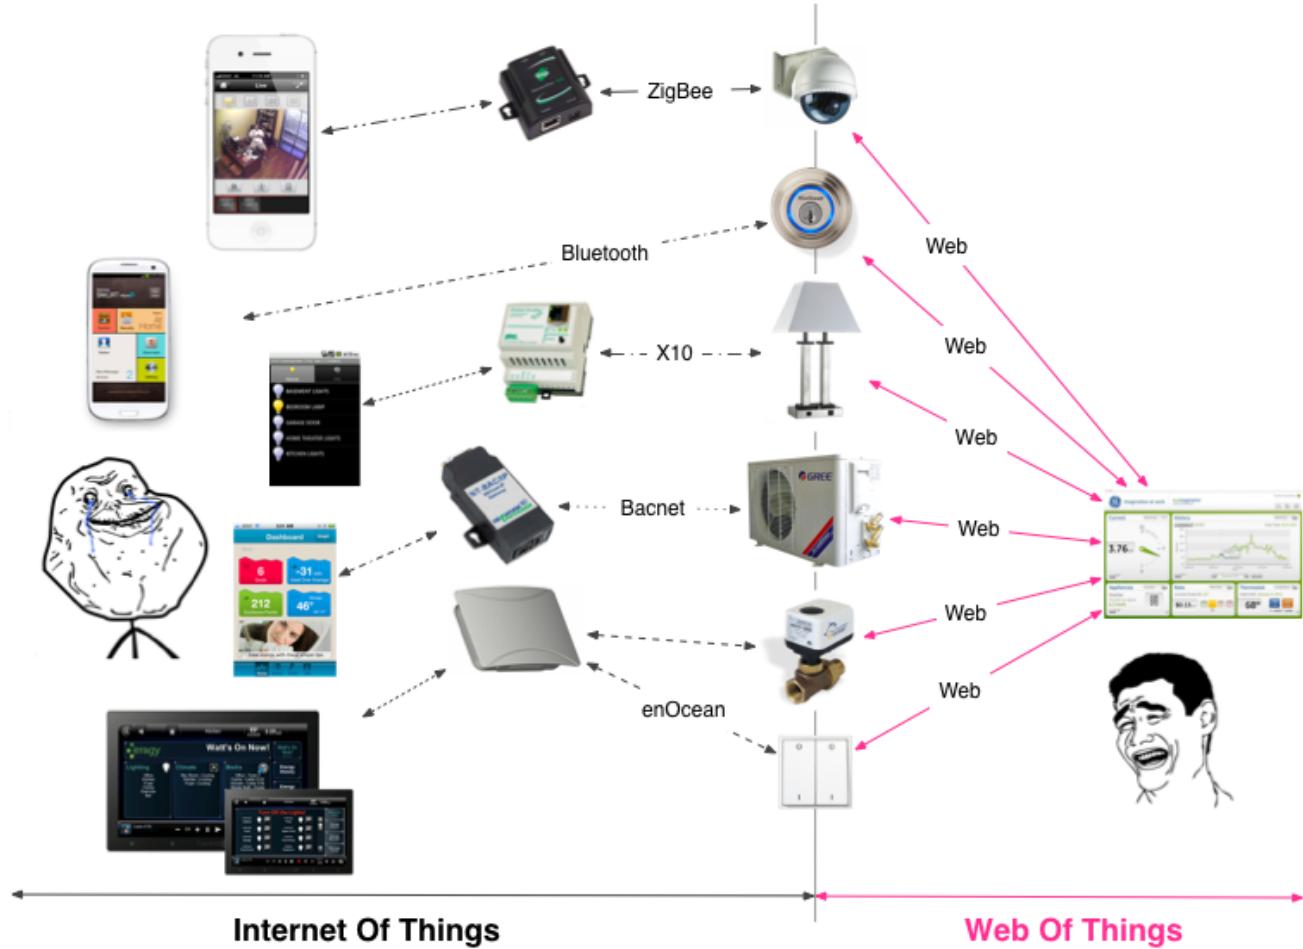
- What is a Web-Service? Why C++?
- Evolution of a real-time bidding service
- Plot twist: JSON love & hate
- Getting nerdy
- Sprinkled with some benchmarks

¿Web-Service?



<http://born2invest.com/cdn/wp-content/uploads/2015/04/21.png>



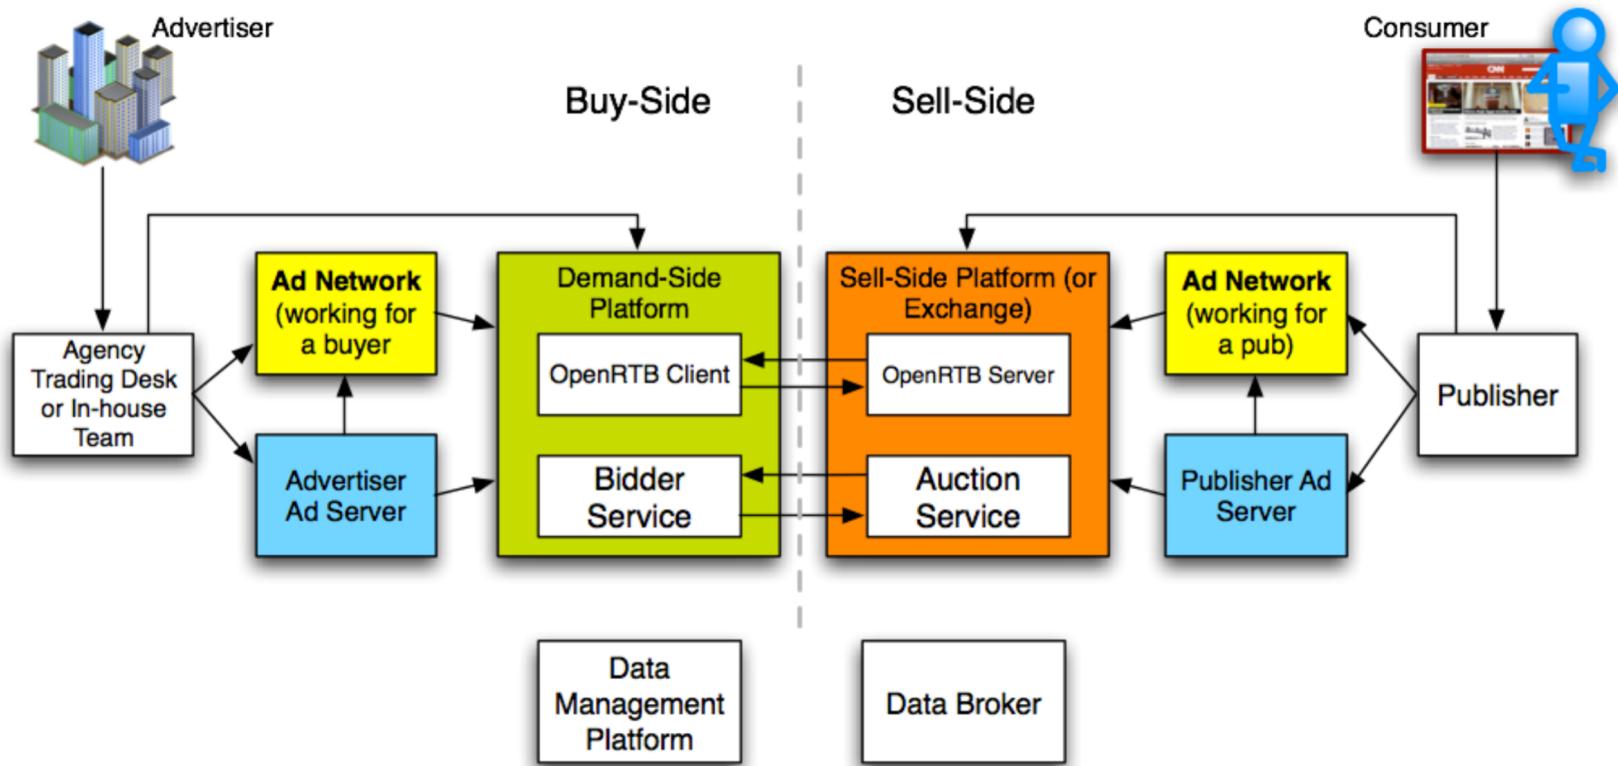


<https://dzone.com/storage/temp/357042-wot-iot.png>

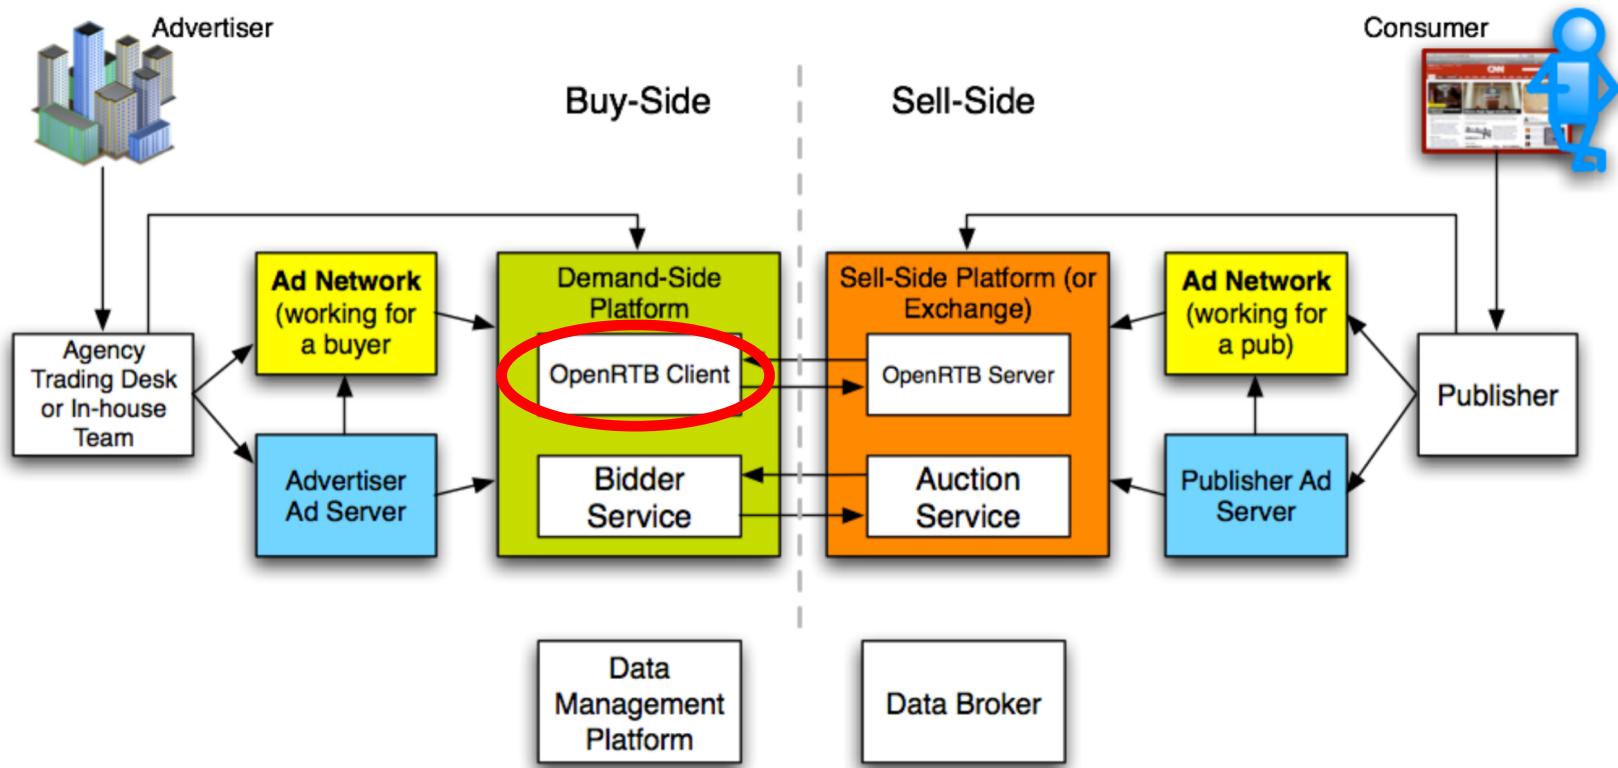
Why Web-Services in C++?

- Response time of single request (e.g. AWS Lambda)
- Flexibility, e.g. shared memory
- Scarce resources
- Many great libraries (TensorFlow, Vowpal Wabbit, OpenCV, ...)
- Existing code-base
- Extensive and mature tools
- C11 and C++11 are beautiful languages ;-)

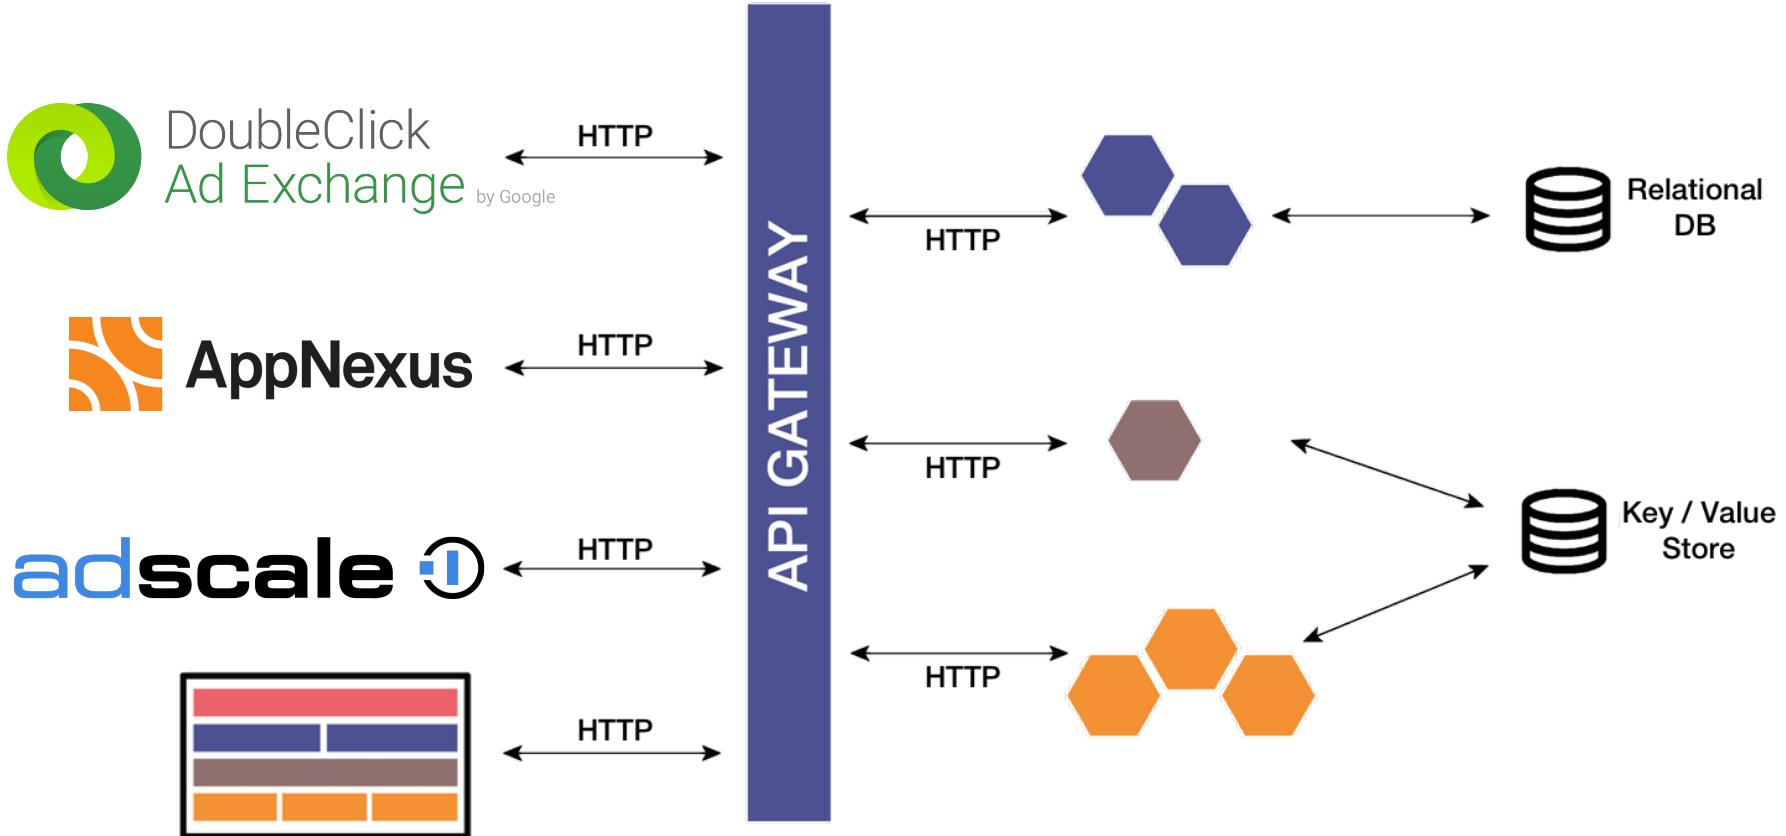
Real-time bidding



Real-time bidding



mbrtargeting's real-time stack (kind of...)



http + json + /bid

HTTP

```
POST /bid HTTP/1.1
Host: localhost:8080
Content-Type: application/json
Content-Length: 13

{"id": "4223"}
```

```
HTTP/1.1 200 OK
Content-Length: 12
Content-Type: text/plain
```

Hello World!

HTTP + X

- http-parser: github.com/nodejs/http-parser
- httppp: github.com/daedric/httppp
- lwan: github.com/lpereira/lwan
- libevhttp: github.com/ellzey/libevhttp
- CivetWeb: github.com/civetweb/civetweb
- Kclone: github.com/koanlogic/klone
- Proxygen: github.com/facebook/proxygen
- and many more: Mongoose, Mongrel2, libwebsockets, QtWebApp, Wt, ...

libevhttp hello world

```
#include <evhttp.h>

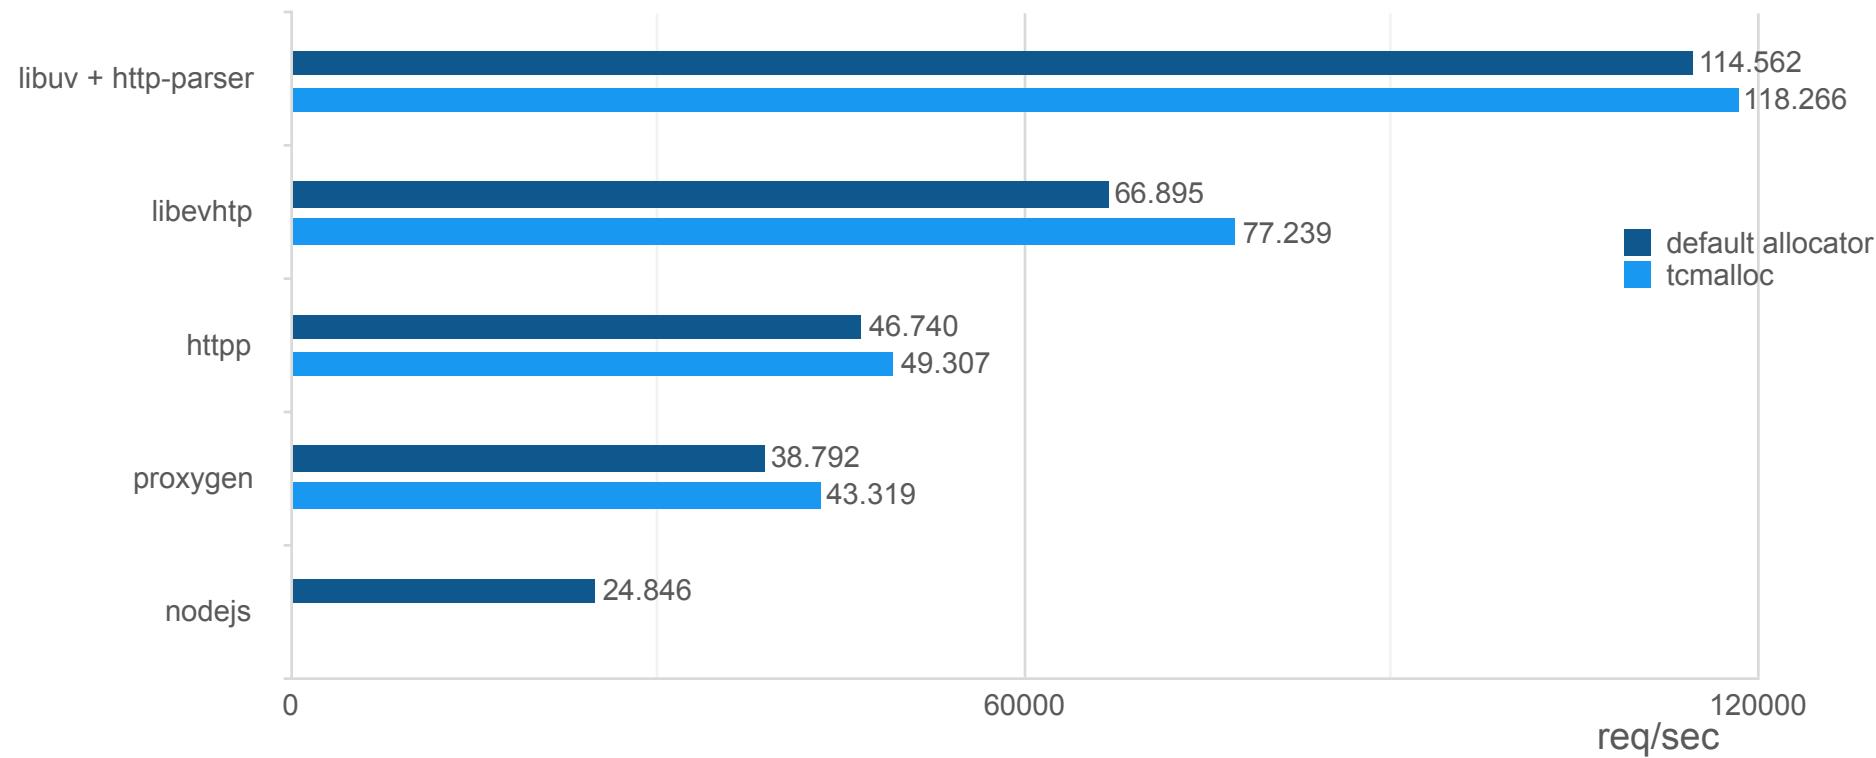
void hellocb(evhttp_request_t* req, void* data) {
    evbuffer_add_printf(req->buffer_out, "Hello World!");
    evhttp_send_reply(req, EVHTTP_RES_OK);
}

int main() {
    evbase_t* evbase = event_base_new();
    evhttp_t* http = evhttp_new(evbase, NULL);
    evhttp_set_cb(http, "/hello", hellocb, NULL);
    evhttp_use_threads_wexit(http, NULL, NULL, 4, NULL);
    evhttp_bind_socket(http, "0.0.0.0", 8080, 2048);
    event_base_loop(evbase, 0);
}
```

Benchmark interlude

- Never trust other people's benchmarks
- Highly recommend Brendan Gregg's talks and blog: brendangregg.com
- All code is online: github.com/0x7f/cpp-meetup
- MacBook Pro i7 2.3 GHz; wrk (c=100, t=3) POST [bidrequest.2.json](#) (~1.5kb)
- Also see abuv (github.com/0x7f/abuv), libuv-based ab drop-in replacement

Benchmark #1 – raw HTTP

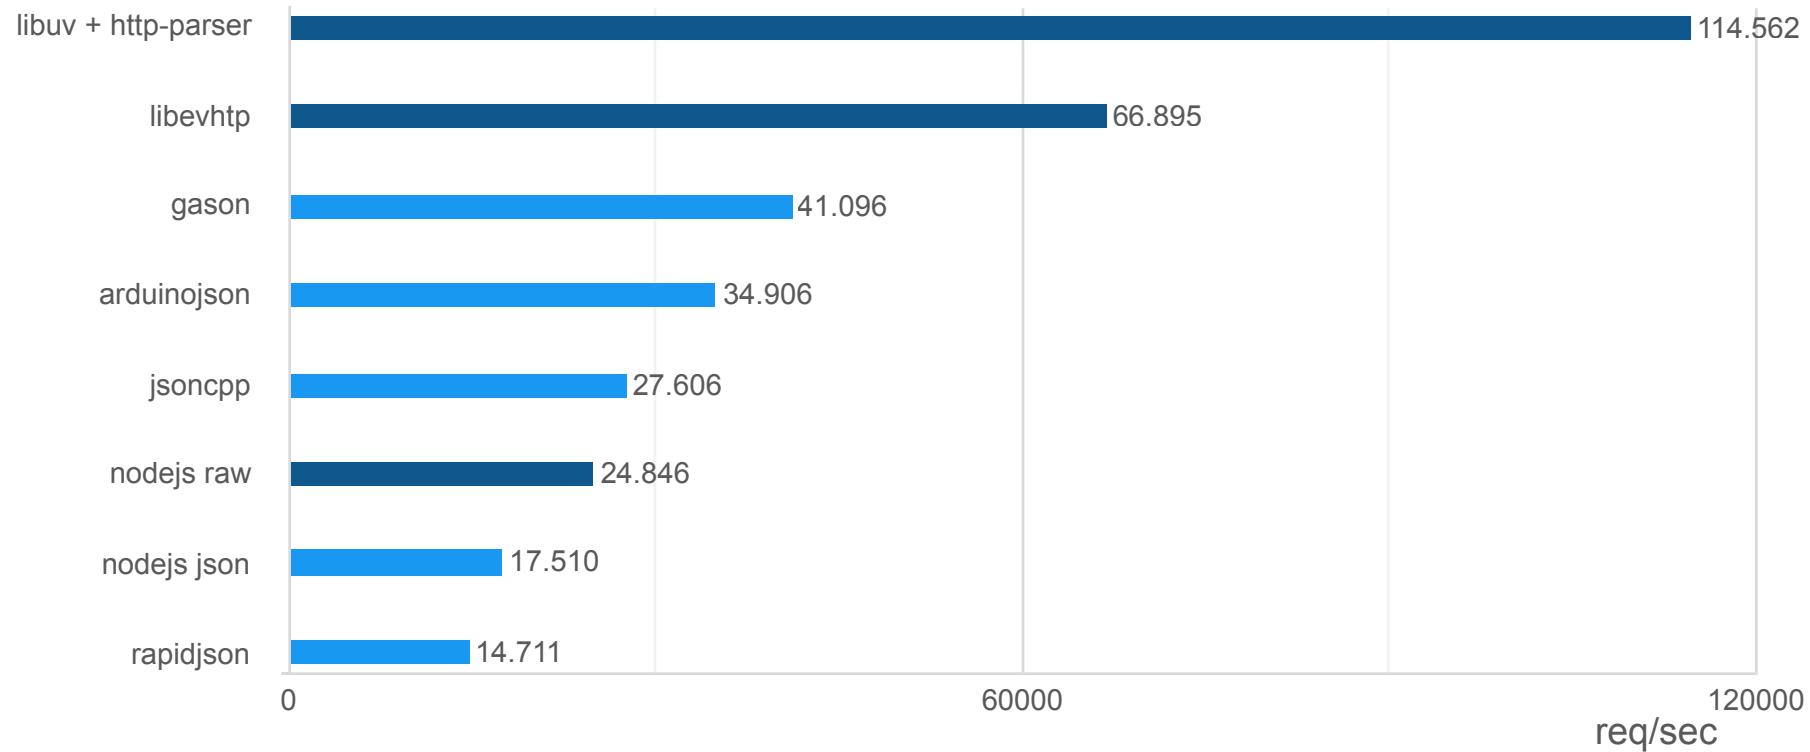


JSON

- JsonCpp github.com/open-source-parsers/jsoncpp
- RapidJSON github.com/miloyip/rapidjson
- Json11 github.com/dropbox/json11
- yajl lloyd.github.io/yajl/
- gason github.com/vivkin/gason
- ArduinoJson github.com/bblanchon/ArduinoJson

For a more complete list, e.g.: github.com/miloyip/nativejson-benchmark

Benchmark #2 – HTTP + JSON



There is no silver bullet

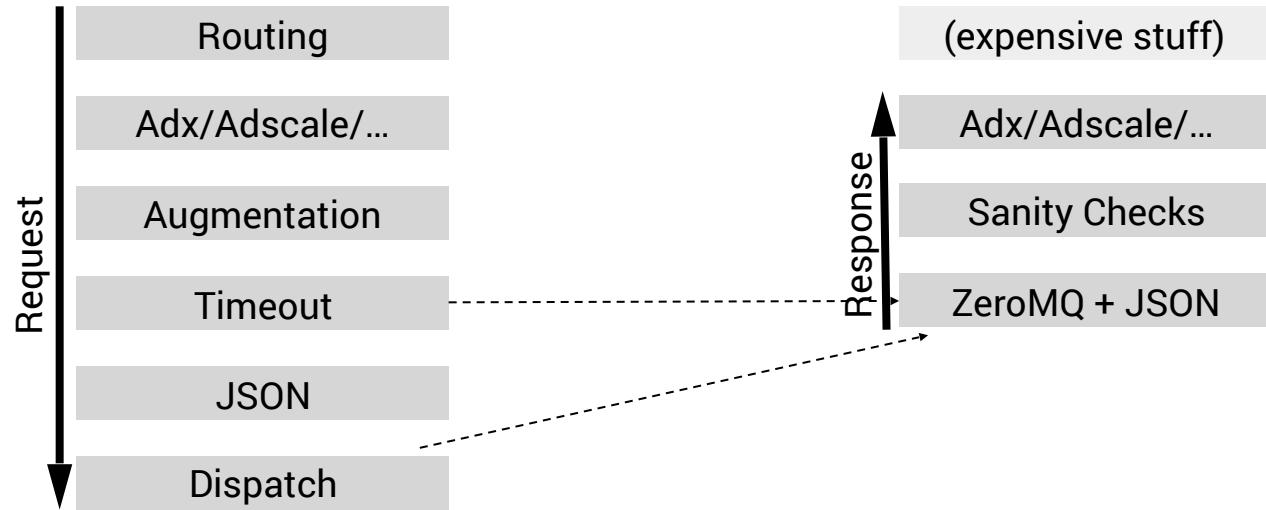
- Green field vs existing code
- Standalone vs embedded http server
- Platform & License requirements
- Multi-purpose backend vs single task
- Serialization vs meat
- Normal vs high throughput
- Hard memory constraints



<https://lisavandore.com/2015/04/17/8-popular-countertop-materials-the-pros-and-the-cons/>

TODO: add some nice graphs of broker architecture

Linux, C++14 (g++ 4.8), boost::asio, http, JsonCpp, ZeroMQ, tcmalloc, Intel TBB



Btw: use latest JsonCpp and link it statically!



Broker evolution

- Augmentation with additional data
- Logging to Kafka
- QoS, load balancing, rate limiting
- Data scientists love data, i.e. traffic

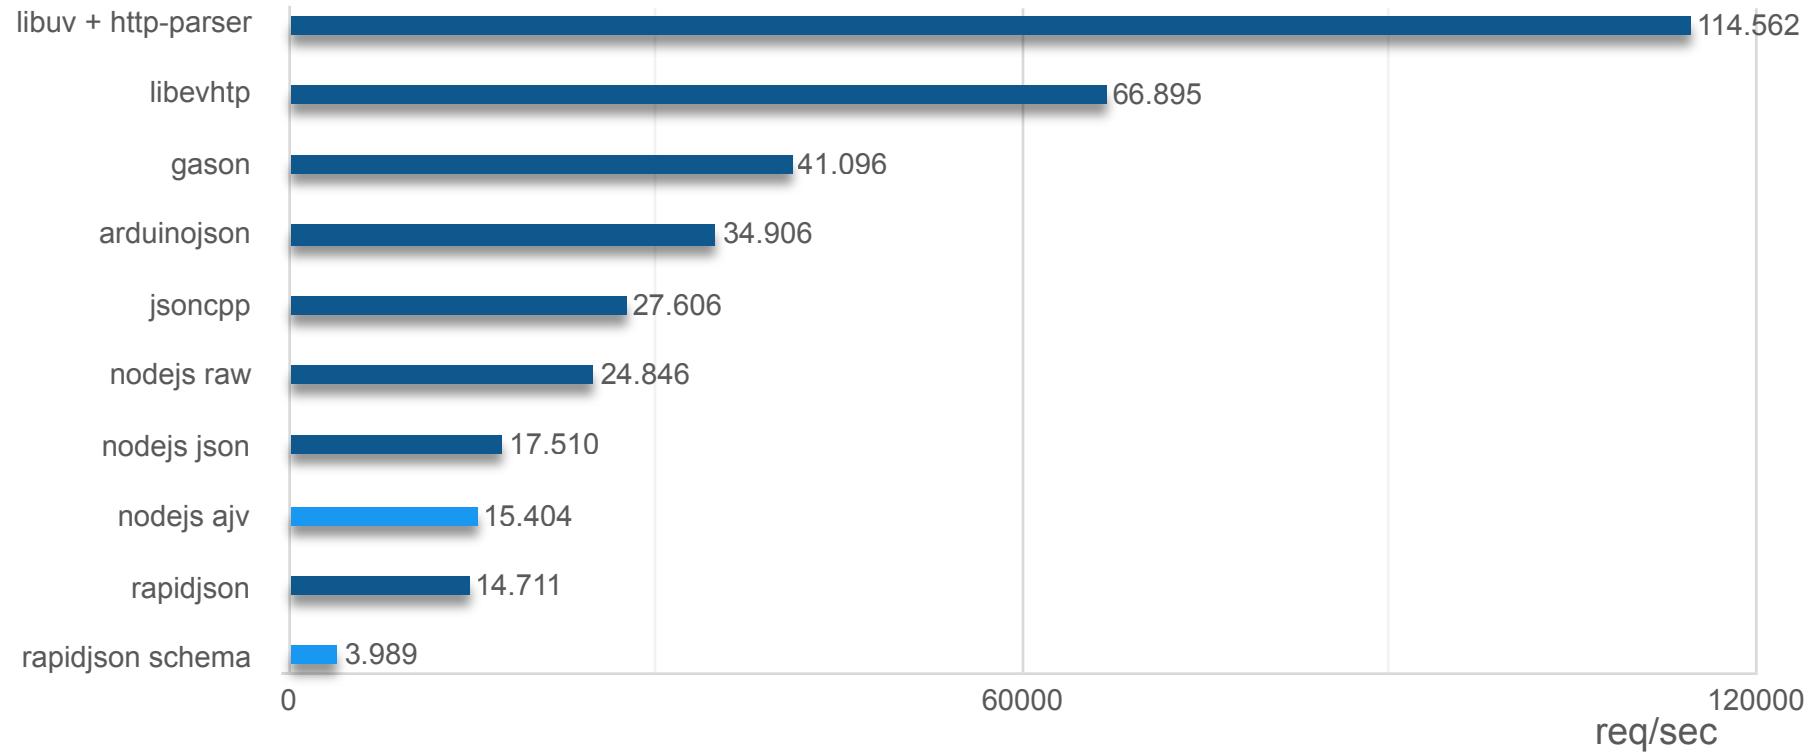
JSON love & hate

```
void doSomething(Json::Value& json) {
    const auto& v = json["myfield"];
    if (!v.isNull() && !v.isString()) {
        throw InvalidBidRequestFormatException();
    }
    if (v.isNull()) {
        return ...;
    }
    const auto i = v.asUInt64();
    ...
}
```

Schemas to the rescue

- JsonSchema helps. a bit. maybe
 - Typos are still an issue
 - Knowledge about types still spread throughout the code
 - Still accessing JSON members through strings
- Start writing your own serialization
 - Structs plus hand-crafted de-/serialization code

Benchmark #3 – HTTP + JSON + Schema



Protobuf to the rescue!!!11elf

```
option java_outer_classname = "Example";
package io. echtzeit.example;

message MsgA {
    optional int32 a = 1;
    required MsgB b = 2;
}

message MsgB {
    required string c = 1;
}
```

Protobuf to the rescue!!!elf – generated Classes

```
option java_outer_classname = "Example";
package io. echtzeit.example;

message MsgA {
    optional int32 a = 1;
    required MsgB b = 2;
}

message MsgB {
    required string c = 1;
}

// optional int32 a = 1;
inline bool has_a() const;
inline void clear_a();
static const int kAFieldNumber = 1;
inline ::google::protobuf::int32 a() const;
inline void set_a(::google::protobuf::int32 value);

// required .io. echtzeit.example.MsgB b = 2;
inline bool has_b() const;
inline void clear_b();
static const int kBFieldNumber = 2;
inline const ::io:: echtzeit::example::MsgB& b() const;
inline ::io:: echtzeit::example::MsgB* mutable_b();
inline ::io:: echtzeit::example::MsgB* release_b();
inline void set_allocated_b(::io:: echtzeit::example::MsgB* b);
```

Protobuf to the rescue!!!elf – Reflection

```
Message* msg = new MsgA();
const Descriptor* descriptor = msg->GetDescriptor();

const FieldDescriptor* numbers_field = descriptor->FindFieldByName( "a" );
assert(numbers_field != NULL);
assert(numbers_field->type() == FieldDescriptor::TYPE_INT32);

// Parse the message.
msg->ParseFromString(data);

const Reflection* reflection = msg->GetReflection();
assert(reflection->GetInt32(msg, numbers_field) == 1);
```

Protobuf to the rescue!!!elf – Serialization

```
// Parse a protocol buffer from a file descriptor.  If successful, the entire
// input will be consumed.
bool ParseFromDescriptor(int file_descriptor);
// Like ParseFromDescriptor(), but accepts messages that are missing
// required fields.
bool ParsePartialFromDescriptor(int file_descriptor);
// Parse a protocol buffer from a C++ istream.  If successful, the entire
// input will be consumed.
bool ParseFromIstream(istream* input);
// Like ParseFromIstream(), but accepts messages that are missing
// required fields.
bool ParsePartialFromIstream(istream* input);

// Serialize the message and write it to the given file descriptor.  All
// required fields must be set.
bool SerializeToFileDescriptor(int file_descriptor) const;
// Like SerializeToFileDescriptor(), but allows missing required fields.
bool SerializePartialToFileDescriptor(int file_descriptor) const;
// Serialize the message and write it to the given C++ ostream.  All
// required fields must be set.
bool SerializeToOstream(ostream* output) const;
// Like SerializeToOstream(), but allows missing required fields.
bool SerializePartialToOstream(ostream* output) const;
```

Protobuf to the rescue!!!elf – Message::Clear()

```
MsgA* New() const;
void CopyFrom(const ::google::protobuf::Message& from);
void MergeFrom(const ::google::protobuf::Message& from);
void CopyFrom(const MsgA& from);
void MergeFrom(const MsgA& from);
void Clear();
bool IsInitialized() const;
```

Protobuf to the rescue!!!elf – Documentation

- Thanks Google: [openrtb.proto \(github.com/google/openrtb\)](https://github.com/google/openrtb)

```
message BidRequest {
    // Unique ID of the bid request, provided by the exchange.
    // REQUIRED by the OpenRTB specification.
    // [AdX: BidRequest.id]
    required string id = 1;

    // Array of Imp objects (Section 3.2.2) representing the impressions offered.
    // At least 1 Imp object is required.
    // [AdX: BidRequest.AdSlot]
    repeated Imp imp = 2;

    oneof distributionchannel_oneof {
```

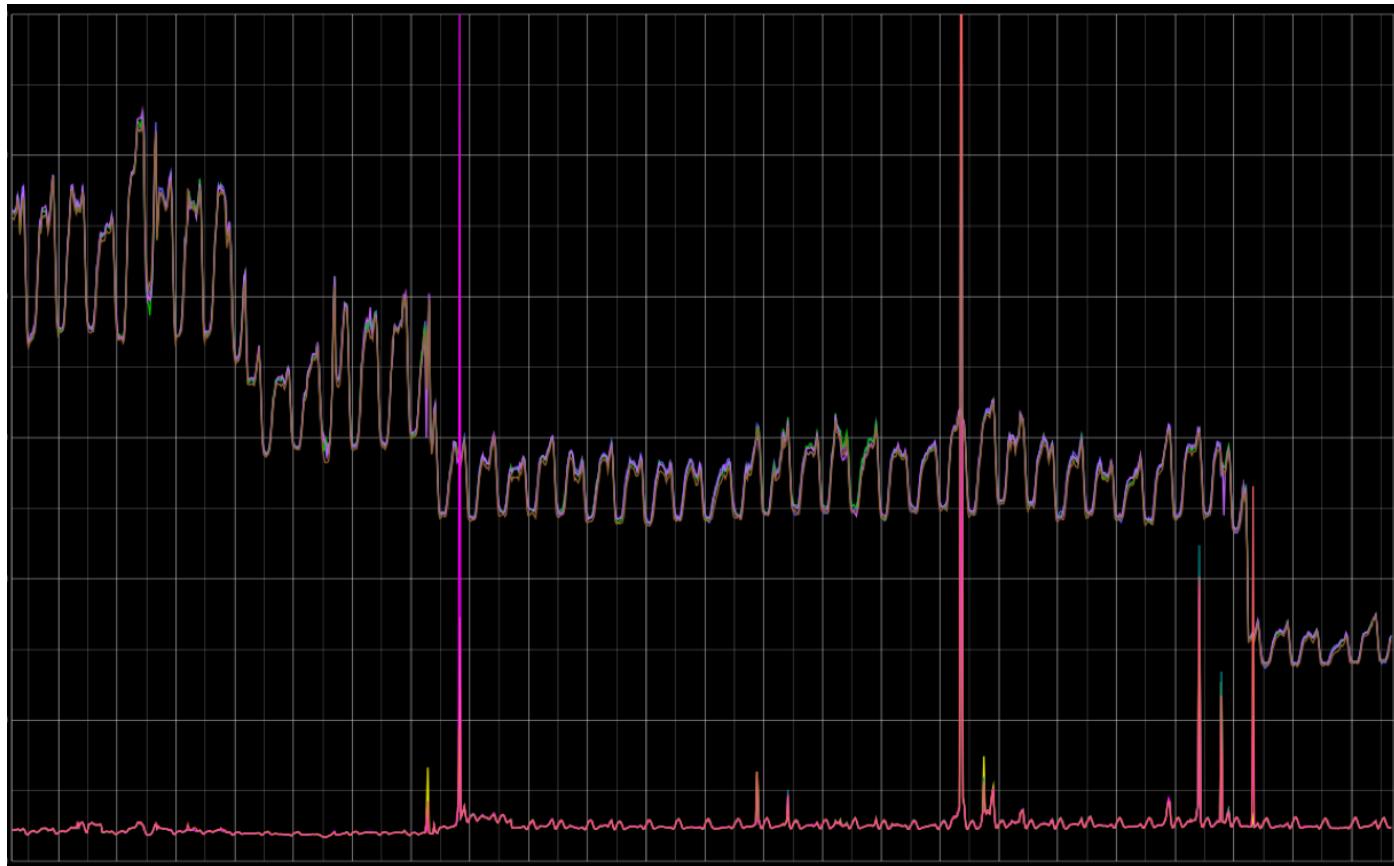
JSON to Protobuf and back

- Google Protocol Buffers 3: [google/protobuf/util/json_util.h](https://github.com/google/protobuf/util/json_util.h)
 - json-protobuf mapper using the generated reflection
 - github.com/yinqiwen/pbjson (RapidJSON based)
 - github.com/shramov/json2pb (jansson based)
 - github.com/shafreeck/pb2json (jansson based)
 - github.com/bivas/protobuf-java-format (Java)
 - github.com/mafintosh/protocol-buffers (node.js)
- Single place for schema validation and performance improvements

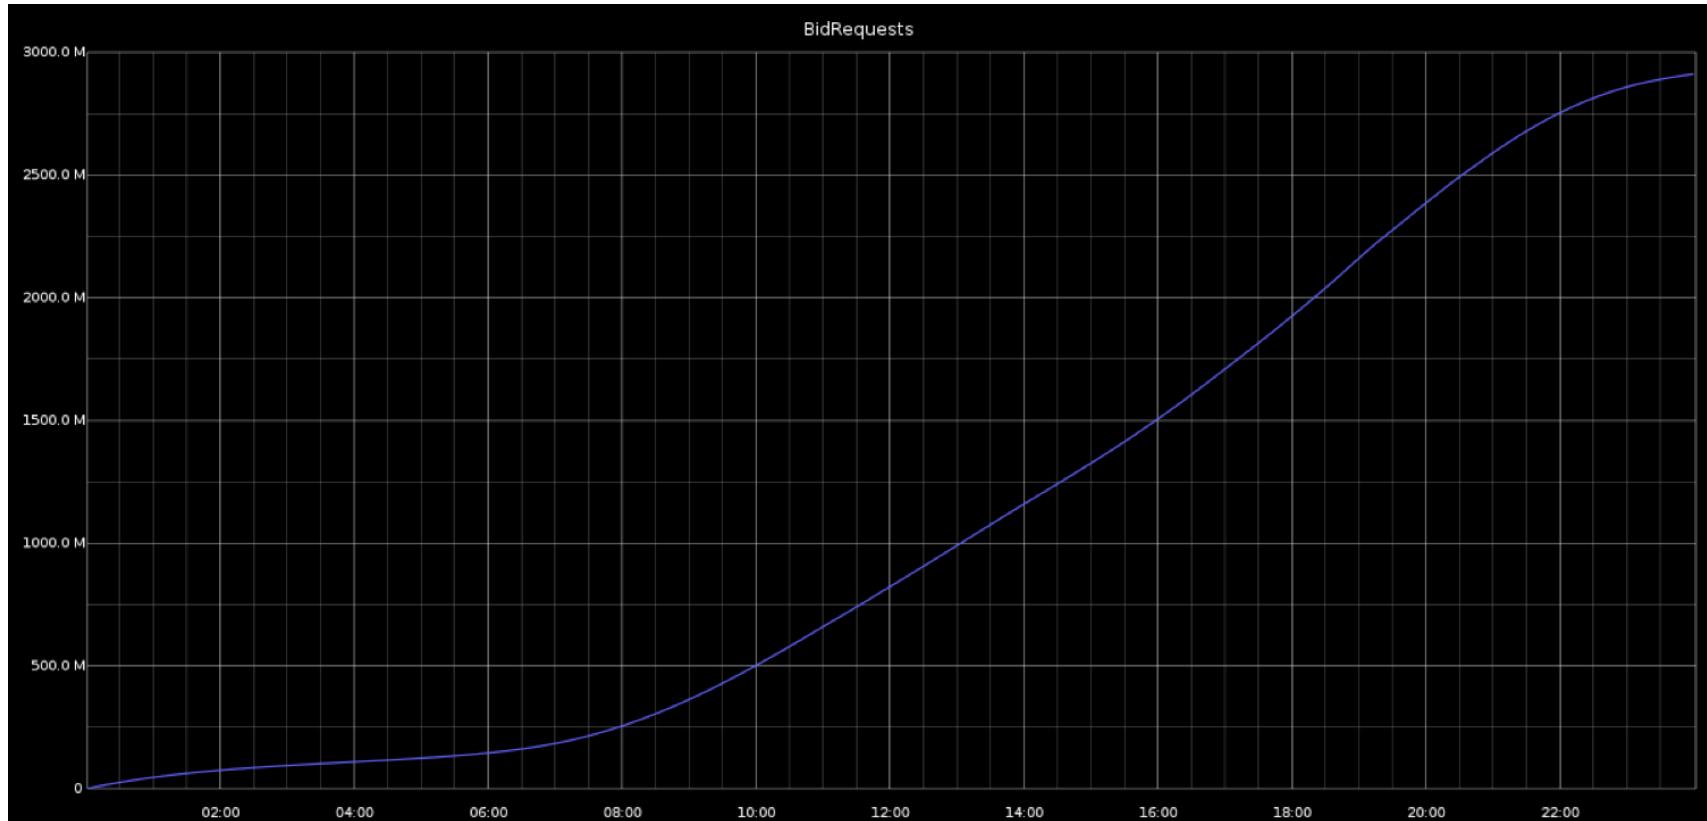
JSON to Protobuf and back

```
BidRequest bidRequest;
auto result = JsonStringToMessage(body, &bidRequest);
if (result.ok()) {
    auto greeting = "Hello " + bidRequest.id();
    response.setCode(HttpStatusCode::Ok).setBody(greeting);
} else {
    response.setCode(HttpStatusCode::BadRequest)
        .setBody(std::string(result.error_message()));
}
```

Drop the bass response times



Now ~3 billion requests/day



Let it sink

...

Use Protobuf for all data structures!



Sky is the limit

- Still more memory used than needed (raw buffer + JSON + Protobuf)
- Use event-based JSON parser (e.g. [yajl](#) or [RapidJson](#)) for memory-efficiency
- Get rid of reflection (e.g. by also generating parser code)

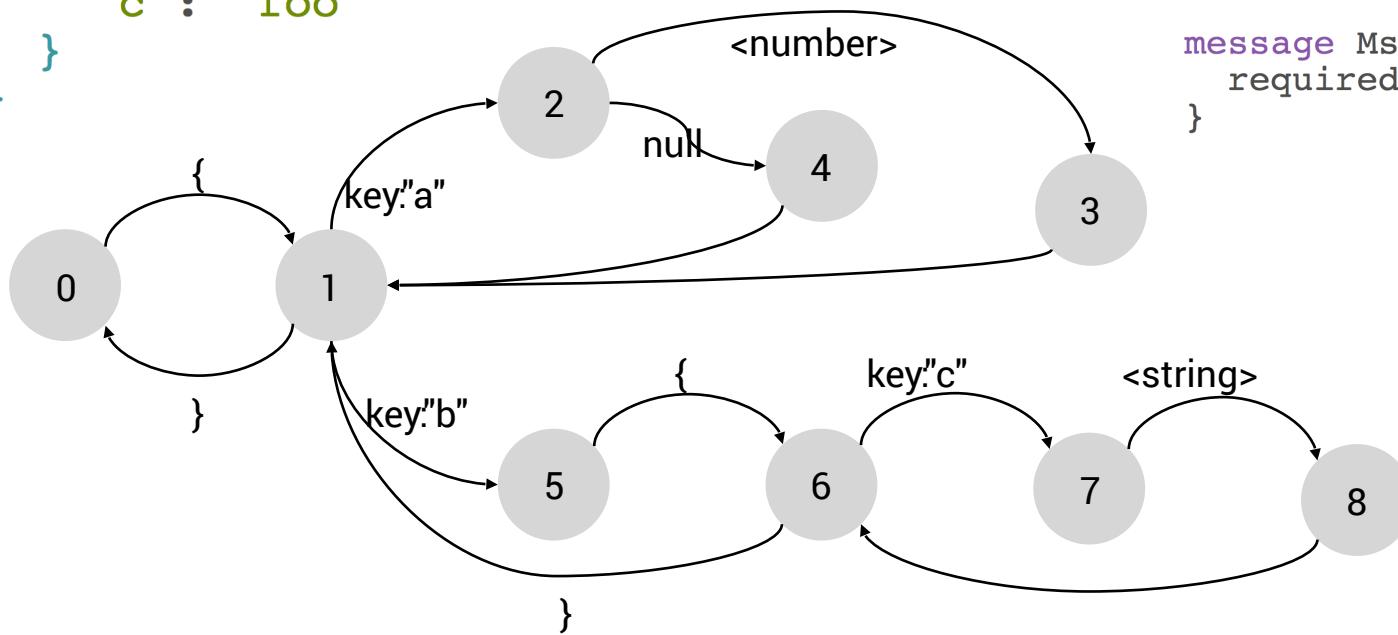
Generate JSON state machine from proto file

```
{  
  "a": 100,  
  "b": {  
    "c": "foo"  
  }  
}
```

```
message MsgA {  
  optional int32 a = 1;  
  required MsgB b = 2;  
}  
  
message MsgB {  
  required string c = 1;  
}
```

Generate JSON state machine from proto file

```
{  
  "a": 100,  
  "b": {  
    "c": "foo"  
  }  
}
```



```
message MsgA {  
  optional int32 a = 1;  
  required MsgB b = 2;  
}
```

```
message MsgB {  
  required string c = 1;  
}
```

Proof-of-concept: protog

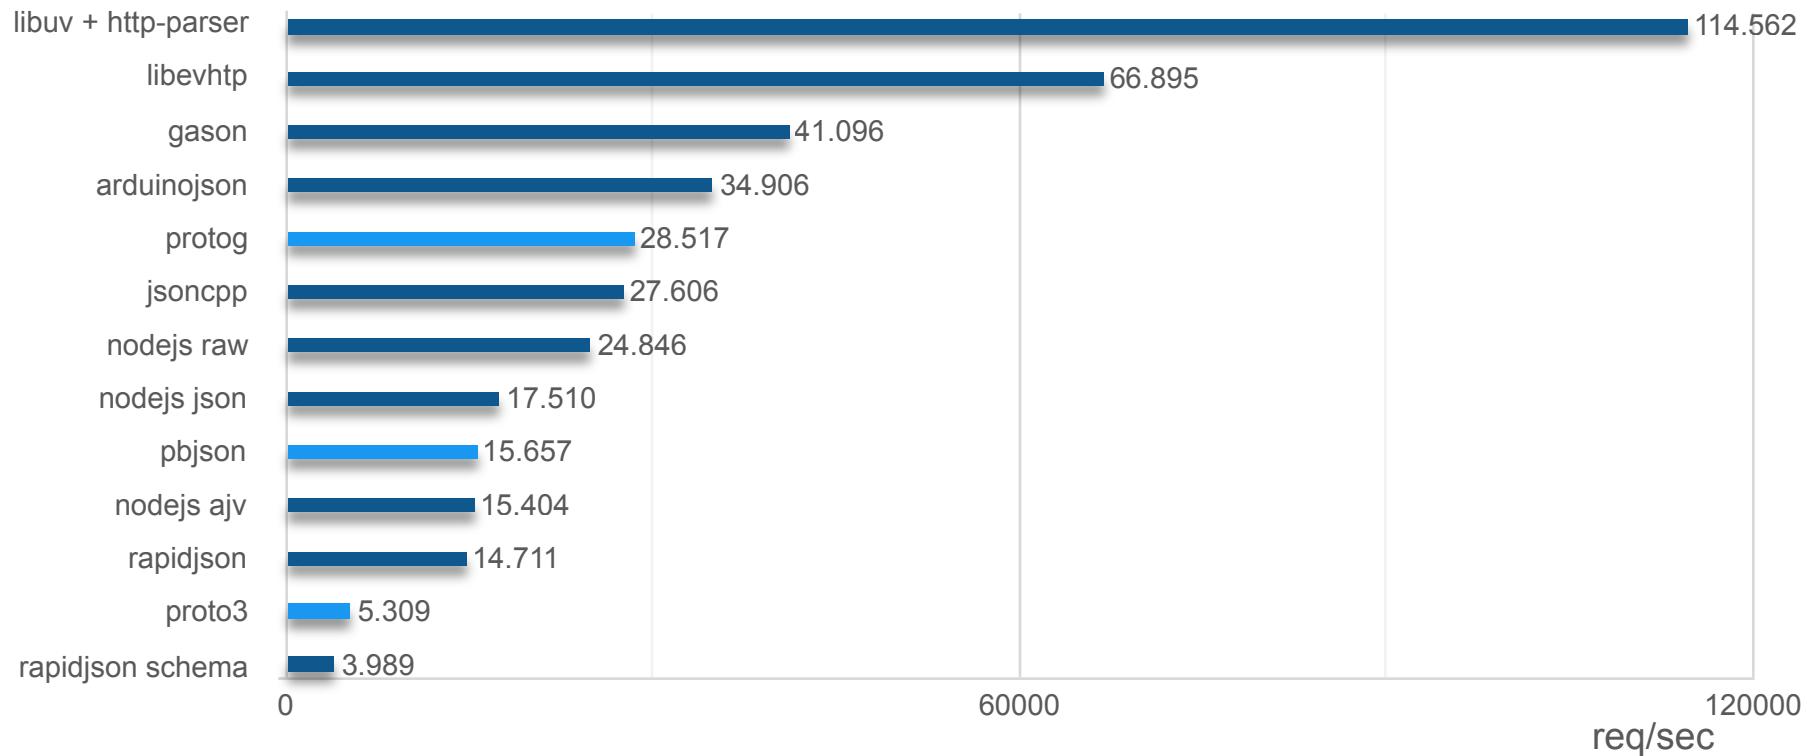
A protobuf json-parser generator!

<https://github.com/0x7f/protog>

```
$ protoc --cpp_out=. openrtb.proto
$ protog -p openrtb.proto -i openrtb.pb.h \
    -m com.google.openrtb.BidRequest
```

Will generate openrtb.pb.{cc,h} and bidrequest_parser.pb.{cc,h}

Benchmark #4 – HTTP + JSON + Protobuf



Summary

- Measure! Measure! Measure!
- Web-services are more than HTTP+JSON
- Mastering JSON is hard
- C++ is not fast by default
- But can be very fast when used correctly
- No silver bullet – heavily depends on the use-case

Questions?



Maximilian Haupt
mhaupt@echtzeitstudios.com
github.com/0x7f